

Lecture 11: Policy Gradients and Actor Critics

Hado van Hasselt

UCL, February 25, 2020

Background reading: Sutton & Barto, 2018, Chapter 13

Vapnik's rule

"Never solve a more general problem as an intermediate step."

— Vladimir Vapnik, 1998

If we care about optimal behaviour: why not learn a policy directly?

General overview

► Model-based RL

- + 'Easy' to learn a model (supervised learning)
- + Learns 'all there is to know' from the data
 - Uses compute & capacity on irrelevant details
 - Computing policy (=planning) is non-trivial and expensive (in compute)

► Value-based RL

- + Easy to generate policy (e.g., $\pi(a|s) = \mathcal{I}(a = \underset{a}{\operatorname{argmax}} q(s, a))$)
- + Close to true objective
- + Fairly well-understood, good algorithms exist
 - Still not the true objective:
 - May focus capacity on irrelevant details
 - Small value error can lead to larger policy error

► Policy-based RL

- + Right objective!
 - o More pros and cons on later slide

General overview

Model-based RL

Value-based RL

Policy-based RL

- ▶ All of these generalise in different ways
- ▶ Sometimes learning a model is easier (e.g., simple dynamics)
- ▶ Sometimes learning a policy is easier (e.g., “always move forward” is optimal)

Policy-Based Reinforcement Learning

- ▶ Previously we approximated parametric value functions

$$v_{\mathbf{w}}(s) \approx v_{\pi}(s)$$
$$q_{\mathbf{w}}(s, a) \approx q_{\pi}(s, a)$$

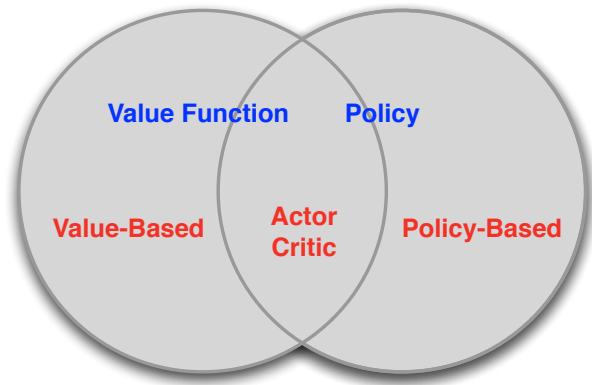
- ▶ A policy can be generated from these values (e.g., greedy)
- ▶ In this lecture we directly parametrise the **policy** directly

$$\pi_{\theta}(a|s) = p(a|s, \theta)$$

- ▶ This lecture, we focus on **model-free** reinforcement learning

Value-based and policy-based RL: terminology

- ▶ **Value Based**
 - ▶ Learnt values
 - ▶ Implicit policy (e.g. ϵ -greedy)
- ▶ **Policy Based**
 - ▶ No values
 - ▶ Learnt policy
- ▶ **Actor-Critic**
 - ▶ Learnt values
 - ▶ Learnt policy



Advantages and disadvantages of policy-based RL

Advantages:

- ▶ True objective
- ▶ Easy extended to **high-dimensional** or **continuous** action spaces
- ▶ Can learn **stochastic** policies
- ▶ Sometimes policies are **simple** while values and models are complex
 - ▶ E.g., complicated dynamics, but optimal policy is always “move forward”

Disadvantages:

- ▶ Could get stuck in local optima
- ▶ Obtained knowledge can be **specific**, does not always generalise well
- ▶ Does not necessarily extract all useful information from the data (when used in isolation)

Stochastic policies

Why could we need stochastic policies?


- ▶ In MDPs, there is always an optimal **deterministic** policy
- ▶ But, most problems are **not fully observable**
 - ▶ This is the common case, especially with function approximation
- ▶ The optimal policy may then be stochastic
- ▶ Additional benefits:
 - ▶ Search space is smoother for stochastic policies
⇒ we can use gradients
 - ▶ Provides easy 'exploration' during learning

Example: Rock-Paper-Scissors



- ▶ Two-player game of rock-paper-scissors
 - ▶ Scissors beats paper
 - ▶ Rock beats scissors
 - ▶ Paper beats rock
- ▶ Consider **iterated** rock-paper-scissors
 - ▶ A deterministic policy is easily exploited
 - ▶ A uniform random policy is optimal (i.e. Nash equilibrium)

Example: Aliased Gridworld




- ▶ The grey states look the same
- ▶ Consider features:

$$\phi(s) = \overbrace{\left(\underbrace{1}_{\text{up}} \underbrace{0}_{\text{right}} \underbrace{1}_{\text{down}} \underbrace{0}_{\text{left}} \right)}^{\text{walls=state}}$$


- ▶ Compare **deterministic** and **stochastic** policies

Example: Aliased Gridworld



- ▶ Under aliasing, an optimal **deterministic** policy will either
 - ▶ move left in both grey states (shown by red arrows)
 - ▶ or move right in both grey states
- ▶ Either way, it can get stuck and never reach the money

Example: Aliased Gridworld (3)



- ▶ An optimal **stochastic** policy moves randomly E or W in grey states

$$\pi_{\theta}(\text{move right} \mid \text{wall up and down}) = 0.5$$

$$\pi_{\theta}(\text{move left} \mid \text{wall up and down}) = 0.5$$

- ▶ Will reach the goal state in a few steps with high probability
- ▶ Policy-based RL can learn the optimal stochastic policy
- ▶ Also when optimal policy does not give equal probability
(So this differs from random tie-breaking with values.)

Policy Objective Functions

- ▶ Goal: given **policy** $\pi_{\theta}(s, a)$, find best **parameters** θ
- ▶ How do we measure the quality of a policy π_{θ} ?
- ▶ In episodic environments we can use the **average total return per episode**
- ▶ In continuing environments we can use the **average reward per step**

Policy Objective Functions

► Episodic-return objective:

$$J_G(\theta) = \mathbb{E}_{\pi_\theta} \left[\sum_{t=0}^{\infty} \gamma^t R_{t+1} \mid S_0 \sim d_0 \right] = \mathbb{E}_{\pi_\theta} [G_0 \mid S_0 \sim d_0] = \mathbb{E}[v_{\pi_\theta}(S_0) \mid S_0 \sim d_0]$$

where d_0 is the start-state distribution

(This objective equals the expected value of the start state)

► Average-reward objective

$$J_R(\theta) = \mathbb{E}_{\pi_\theta} [R] = \sum_s d_{\pi_\theta}(s) \sum_a \pi_\theta(s, a) \sum_r p(r \mid s, a) r$$

where $d_\pi(s) = p(S_t = s \mid \pi)$ is the probability of being in state s in the long run

Think of it as the ratio of time spent in s under policy π

Policy Optimisation

- ▶ Policy based reinforcement learning is an **optimization** problem
- ▶ Find θ that maximises $J(\theta)$
- ▶ We will focus on **stochastic gradient ascent**, which is often quite efficient (and easy to use with deep nets)
- ▶ Some approaches do not use gradient
 - ▶ Hill climbing / simulated annealing
 - ▶ Genetic algorithms / evolutionary strategies

Policy Gradient



- Idea: ascent the gradient of the objective $J(\theta)$

$$\Delta\theta = \alpha \nabla_{\theta} J(\theta)$$

- Where $\nabla_{\theta} J(\theta)$ is the **policy gradient**

$$\nabla_{\theta} J(\theta) = \begin{pmatrix} \frac{\partial J(\theta)}{\partial \theta_1} \\ \vdots \\ \frac{\partial J(\theta)}{\partial \theta_n} \end{pmatrix}$$

- and α is a step-size parameter
- Stochastic policies help ensure $J(\theta)$ is smooth (typically/mostly)



Gradients on parameterized policies

- ▶ How to compute this gradient $\nabla_{\theta} J(\theta)$?
- ▶ Assume policy π_{θ} is differentiable almost everywhere (e.g., neural net)
- ▶ For average reward

$$\nabla_{\theta} J(\theta) = \nabla_{\theta} \mathbb{E}_{\pi_{\theta}}[R].$$

- ▶ How does $\mathbb{E}[R]$ depend on θ ?

Contextual Bandits Policy Gradient

- ▶ Consider a one-step case (a contextual bandit) such that $J(\theta) = \mathbb{E}_{\pi_\theta}[R(S, A)]$.
(Expectation is over d (states) and π (actions))
(For now, d does **not** depend on π)
- ▶ We cannot sample R_{t+1} and then take a gradient:
 R_{t+1} is just a number and does not depend on θ !
- ▶ Instead, we use the identity:

$$\nabla_{\theta} \mathbb{E}_{\pi_{\theta}}[R(S, A)] = \mathbb{E}_{\pi_{\theta}}[R(S, A) \nabla_{\theta} \log \pi(A|S)].$$

(Proof on next slide)

- ▶ The right-hand side gives an expected gradient that can be sampled
- ▶ Also known as REINFORCE (Williams, 1992)

The score function trick

Let $r_{sa} = \mathbb{E}[R(S, A) \mid S = s, A = s]$

$$\begin{aligned}\nabla_{\theta} \mathbb{E}_{\pi_{\theta}}[R(S, A)] &= \nabla_{\theta} \sum_s d(s) \sum_a \pi_{\theta}(a|s) r_{sa} \\&= \sum_s d(s) \sum_a r_{sa} \nabla_{\theta} \pi_{\theta}(a|s) \\&= \sum_s d(s) \sum_a r_{sa} \pi_{\theta}(a|s) \frac{\nabla_{\theta} \pi_{\theta}(a|s)}{\pi_{\theta}(a|s)} \\&= \sum_s d(s) \sum_a \pi_{\theta}(a|s) r_{sa} \nabla_{\theta} \log \pi_{\theta}(a|s) \\&= \mathbb{E}_{\pi_{\theta}}[R(S, A) \nabla_{\theta} \log \pi_{\theta}(A|S)]\end{aligned}$$

Contextual Bandit Policy Gradient

$$\nabla_{\theta} \mathbb{E}[R(S, A)] = \mathbb{E}[\nabla_{\theta} \log \pi_{\theta}(A|S) R(S, A)] \quad (\text{see previous slide})$$

- ▶ This is something we **can** sample
- ▶ Our stochastic policy-gradient update is then

$$\theta_{t+1} = \theta_t + \alpha R_{t+1} \nabla_{\theta} \log \pi_{\theta_t}(A_t|S_t).$$

- ▶ In expectation, this is the following the actual gradient
- ▶ So this is a pure stochastic gradient algorithm
- ▶ Intuition: increase probability for actions with high rewards

Policy gradients: reduce variance

- Note that, in general

$$\begin{aligned}\mathbb{E}[b \nabla_{\theta} \log \pi(A_t|S_t)] &= \mathbb{E} \left[\sum_a \pi(a|S_t) b \nabla_{\theta} \log \pi(a|S_t) \right] \\ &= \mathbb{E} \left[b \nabla_{\theta} \sum_a \pi(a|S_t) \right] \\ &= \mathbb{E}[b \nabla_{\theta} 1] &= 0\end{aligned}$$

- This is true if b does not depend on the action (but it can depend on the state)
- Implies we can subtract a **baseline** to reduce variance

$$\theta_{t+1} = \theta_t + \alpha(R_{t+1} - b(S_t)) \nabla_{\theta} \log \pi_{\theta_t}(A_t|S_t).$$

- We will also use this fact in proofs below

Example: Softmax Policy

- ▶ Consider a softmax policy on action preferences $h(s, a)$ as an example
- ▶ Probability of action is proportional to exponentiated weight

$$\pi_{\theta}(a|s) = \frac{e^{h(s,a)}}{\sum_b e^{h(s,b)}}$$

- ▶ The gradient of the log probability is

$$\nabla_{\theta} \log \pi_{\theta}(A_t|S_t) = \underbrace{\nabla_{\theta} h(S_t, A_t)}_{\text{gradient of preference}} - \underbrace{\sum_a \pi_{\theta}(a|S_t) \nabla_{\theta} h(S_t, a)}_{\text{expected gradient of preference}}$$

Policy Gradient Theorem

- ▶ The policy gradient approach also applies to (multi-step) MDPs
- ▶ Replaces reward R with long-term return G_t or value $q_\pi(s, a)$
- ▶ There are actually two policy gradient theorems (Sutton et al., 2000):
average return per episode & **average reward per step**

Policy gradient theorem (episodic)

Theorem

For any differentiable policy $\pi_{\theta}(s, a)$, let d_0 be the starting distribution over states in which we begin an episode. Then, the policy gradient of $J(\theta) = \mathbb{E}[G_0 \mid S_0 \sim d_0]$ is

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\pi_{\theta}} \left[\sum_{t=0}^T \gamma^t \mathbf{q}_{\pi_{\theta}}(S_t, A_t) \nabla_{\theta} \log \pi_{\theta}(A_t | S_t) \mid S_0 \sim d_0 \right]$$

where

$$\begin{aligned} q_{\pi}(s, a) &= \mathbb{E}_{\pi}[G_t \mid S_t = s, A_t = a] \\ &= \mathbb{E}_{\pi}[R_{t+1} + \gamma q_{\pi}(S_{t+1}, A_{t+1}) \mid S_t = s, A_t = a] \end{aligned}$$

Policy gradients on trajectories

- ▶ Policy gradients do **not** need to know the dynamics
- ▶ Kind of surprising; shouldn't we know how the policy influences the states?

Episodic policy gradients: proof

- Consider trajectory $\tau = S_0, A_0, R_1, S_1, A_1, R_1, S_2, \dots$ with return $G(\tau)$

$$\nabla_{\theta} J_{\theta}(\pi) = \nabla_{\theta} \mathbb{E}[G(\tau)] = \mathbb{E}[G(\tau) \nabla_{\theta} \log p(\tau)] \quad (\text{score function trick})$$

$$\begin{aligned} \nabla_{\theta} \log p(\tau) &= \nabla_{\theta} \log \left[p(S_0) \pi(A_0|S_0) p(S_1|S_0, A_0) \pi(A_1|S_1) \cdots \right] \\ &= \nabla_{\theta} \left[\log p(S_0) + \log \pi(A_0|S_0) + \log p(S_1|S_0, A_0) + \log \pi(A_1|S_1) + \cdots \right] \\ &= \nabla_{\theta} \left[\log \pi(A_0|S_0) + \log \pi(A_1|S_1) + \cdots \right] \end{aligned}$$

So:

$$\nabla_{\theta} J_{\theta}(\pi) = \mathbb{E}_{\pi} [G(\tau) \nabla_{\theta} \sum_{t=0}^T \log \pi(A_t|S_t)]$$

Episodic policy gradients: proof (continued)

$$\begin{aligned}\nabla_{\boldsymbol{\theta}} J_{\boldsymbol{\theta}}(\pi) &= \mathbb{E}_{\pi} \left[G(\tau) \sum_{t=0}^T \nabla_{\boldsymbol{\theta}} \log \pi(A_t | S_t) \right] \\&= \mathbb{E}_{\pi} \left[\sum_{t=0}^T G(\tau) \nabla_{\boldsymbol{\theta}} \log \pi(A_t | S_t) \right] \\&= \mathbb{E}_{\pi} \left[\sum_{t=0}^T \left(\sum_{k=0}^T \gamma^k R_{k+1} \right) \nabla_{\boldsymbol{\theta}} \log \pi(A_t | S_t) \right] \\&= \mathbb{E}_{\pi} \left[\sum_{t=0}^T \left(\sum_{k=\textcolor{red}{t}}^T \gamma^k R_{k+1} \right) \nabla_{\boldsymbol{\theta}} \log \pi(A_t | S_t) \right] \\&= \mathbb{E}_{\pi} \left[\sum_{t=0}^T \left(\gamma^t \sum_{k=t}^T \gamma^{k-t} R_{k+1} \right) \nabla_{\boldsymbol{\theta}} \log \pi(A_t | S_t) \right] \\&= \mathbb{E}_{\pi} \left[\sum_{t=0}^T (\gamma^t G_t) \nabla_{\boldsymbol{\theta}} \log \pi(A_t | S_t) \right] \qquad = \mathbb{E}_{\pi} \left[\sum_{t=0}^T \gamma^t q_{\pi}(S_t, A_t) \nabla_{\boldsymbol{\theta}} \log \pi(A_t | S_t) \right]\end{aligned}$$

Episodic policy gradients algorithm

$$\nabla_{\theta} J_{\theta}(\pi) = \mathbb{E}_{\pi} \left[\sum_{t=0}^T \gamma^t q_{\pi}(S_t, A_t) \nabla_{\theta} \log \pi(A_t | S_t) \right]$$

- ▶ We can sample this, given a whole episode
- ▶ Typically, people pull out the sum, and split up this into separate gradients, e.g.,

$$\Delta \theta_t = \gamma^t G_t \nabla_{\theta} \log \pi(A_t | S_t)$$

such that $\mathbb{E}_{\pi} [\sum_t \Delta \theta_t] = \nabla_{\theta} J_{\theta}(\pi)$

- ▶ Typically, people ignore the γ^t term, use $\Delta \theta_t = G_t \nabla_{\theta} \log \pi(A_t | S_t)$
- ▶ This is actually okay-ish — we just partially pretend on each step that we could have started an episode in that state instead (alternatively, view it as a slightly biased gradient)

Policy gradient theorem (average reward)

Theorem

For any differentiable policy $\pi_{\theta}(s, a)$, the policy gradient of $J(\theta) = \mathbb{E}[R \mid \pi]$ is

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\pi} [q_{\pi_{\theta}}(\mathbf{S}_t, \mathbf{A}_t) \nabla_{\theta} \log \pi_{\theta}(\mathbf{A}_t | \mathbf{S}_t)]$$

where

$$q_{\pi}(s, a) = \mathbb{E}_{\pi} [R_{t+1} - \rho + q_{\pi}(S_{t+1}, A_{t+1}) \mid S_t = s, A_t = a]$$

$$\rho = \mathbb{E}_{\pi} [R_{t+1}] \quad (\text{Note: global average, not conditioned on state or action})$$

(Expectation is over both states and actions)

Policy gradients: reduce variance

- ▶ Recall $\mathbb{E}_{\pi}[b(S_t)\nabla \log \pi(A_t|S_t)] = 0$, for any $b(S_t)$ that does not depend on A_t
- ▶ A common baseline is $v_{\pi}(S_t)$

$$\nabla_{\theta} J_{\theta}(\pi) = \mathbb{E} \left[\sum_{t=0}^{\infty} \gamma^t (q_{\pi}(S_t, A_t) - v_{\pi}(S_t)) \nabla_{\theta} \log \pi(A_t|S_t) \right]$$

- ▶ Typically, we estimate $v_{\mathbf{w}}(s) \approx v_{\pi}(s)$ explicitly, and sample

$$q_{\pi}(S_t, A_t) \approx G_t$$

- ▶ We can minimise variance further by **bootstrapping**, e.g., $G_t = R_{t+1} + \gamma v_{\mathbf{w}}(S_{t+1})$